Event forecasting remains exploratory weak-label evidence

Figure S1 | Revised next-new-onset task; demonstration forecasts are separate from deployed robot success

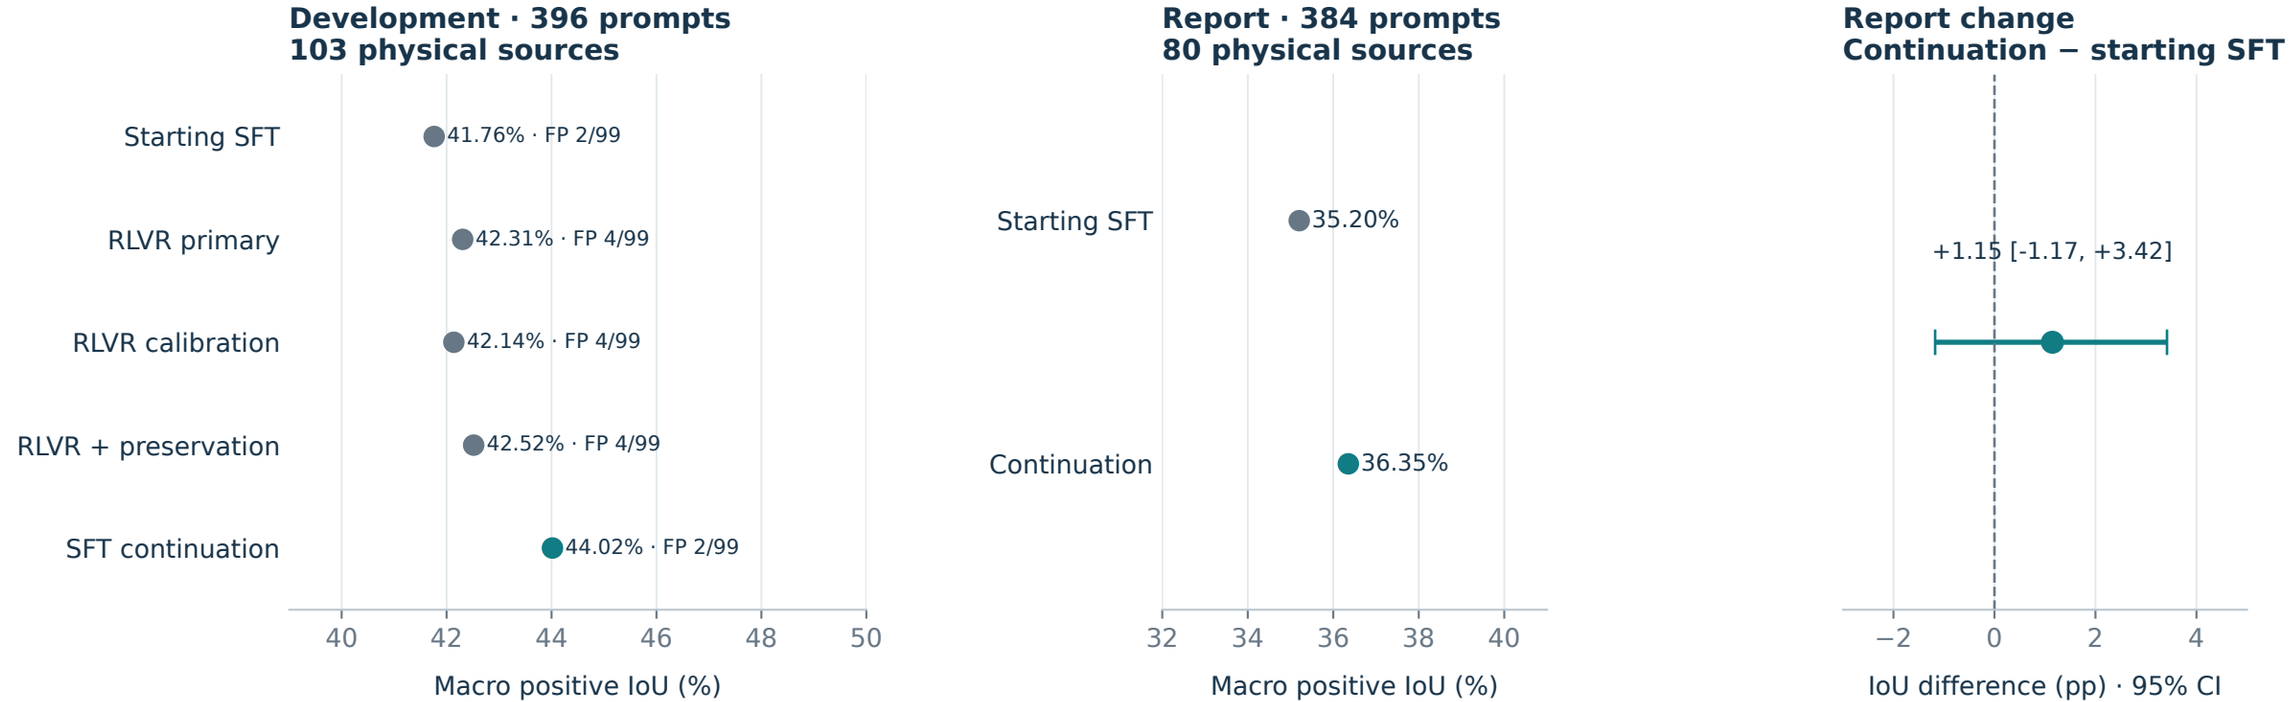

The report interval spans zero. Report none false positives stay at 3/96. SFT continuation was selected on development before report scoring.

No RLVR report arm was added. Automatic targets, adaptive development reuse, prior report exposure and unresolved independent label QA qualify these results.

Report interval: 10,000 paired physical-source draws, PCG64 seed 10. The earlier 48-prompt task has different targets and remains separate.